

ADVANCE DATA MINING AND PREDICTIVE ANALYTICS - 3

shiva gadila

2023-12-06

QA1. What is the difference between SVM with hard margin and soft margin?

In Support Vector Machines (SVM), there are two main approaches: hard margin and soft margin.

Hard Margin SVM: Imagine you have data points of two different classes, and you want to draw a line (hyperplane) that perfectly separates them. This works well when the data is cleanly separable, but in real life, there might be some noisy or outlier points. Hard Margin SVM doesn't handle these outliers gracefully. It insists on a perfect division, and if there's even one misclassified point, it struggles.

Soft Margin SVM: Understanding that perfect separation might not be achievable in the real world, especially with noisy data, the soft margin approach is more forgiving. It introduces a concept called "slack variable." This allows for a bit of flexibility, meaning it permits some points to be on the wrong side of the dividing line. The challenge now becomes finding a balance between having a wide margin (space between classes) and allowing for a few misclassifications. There's a parameter called "C" that decides how much penalty is given for each misclassified point. You can think of it as the cost of allowing mistakes.

Summary: Hard Margin: Insists on a perfect division, not great with noisy data. Soft Margin: More flexible, accepts some misclassifications, and adjusts the balance with the "C" parameter. In simple terms, it's like trying to draw a line between two groups of points. Hard Margin wants a flawless line, while Soft Margin allows for a bit of messiness, deciding how much mess is acceptable with the "C" parameter. It's a way of handling the imperfect nature of real-world data.

QA2. What is the role of the cost parameter, C, in SVM (with soft margin) classification?

The importance of the cost parameter, C, in SVM with a soft margin cannot be overstated. It acts as the guiding force in balancing the priorities of maximizing the margin and minimizing classification errors.

SVM focuses on creating a broader margin to enhance generalization, acknowledging the challenges of achieving perfect separation in real-world scenarios. C allows for some misclassifications but strives to keep them minimal while maintaining a significant margin.

The delicate balance controlled by the C parameter involves a trade-off. Lower C values prioritize a wider margin, accepting a higher tolerance for misclassifications. Conversely, higher C values emphasize reducing misclassifications, potentially leading to a narrower margin.

The impact of the chosen C value is significant. A high C risks overfitting by closely fitting the training data, capturing noise and outliers. In contrast, a low C may result in underfitting as the model leans heavily towards maximizing the margin, potentially neglecting some misclassifications.

Selecting the right C value is crucial for optimal SVM performance. Techniques like cross-validation help explore various C values, ensuring the choice that delivers the best performance on a validation set. This meticulous approach ensures the development of a well-balanced SVM model proficient in managing both margin maximization and classification error minimization.

Q3. Will the following perceptron be activated (2.8 is the activation threshold)?

$$\begin{aligned}\text{Activation function} &= (\text{input 1} \times \text{weight 1}) + (\text{input 2} \times \text{weight 2}) = (0.1 \times 0.8) + (11.1 \times -0.2) \\ &= 0.08 - 2.22 \\ &= -2.14\end{aligned}$$

The activation function value is -2.14, which falls below the activation threshold of 2.8. Consequently, the perceptron will remain inactive in this scenario.

QA4. What is the role of alpha, the learning rate in the delta rule?

In the delta rule, a widely used algorithm for adjusting neural network weights during training, the learning rate, represented by the symbol alpha (α), plays a crucial role. Alpha is a hyperparameter that determines the size of the steps taken for weight updates generated by the delta rule.

Here's how it works: The delta rule calculates the error between the predicted output of the neural network and the actual target output. This error is then utilized to tweak the network's weights, aiming to enhance its overall performance. The extent of weight adjustment is linked to both the magnitude of the error and the learning rate.

Choosing the right value for alpha is key. A higher alpha leads to more substantial weight updates, facilitating faster convergence of weights. However, this can also result in overshooting the optimal weights and hinder convergence. Conversely, a lower alpha results in smaller weight updates and a longer convergence process. While this approach may assist in reaching a more accurate minimum, it requires patience during training.

Selecting an appropriate learning rate is a critical aspect of the delta rule, impacting the algorithm's performance. A common strategy involves starting with a small alpha, such as 0.1 or 0.01, and then fine-tuning it through experimentation to achieve optimal results on the training data.

PART B

#QB1. Build a linear SVM regression model to predict Sales based on all other attributes ("Price", "Advertising", "Population", "Age", "Income" and "Education"). Hint: use caret train() with method set to "svmLinear". What is the R-squared of the model?

```
library(ISLR)
library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(glmnet)

## Loading required package: Matrix
```

```

## Warning: package 'Matrix' was built under R version 4.2.3

## Loaded glmnet 4.1-8

library(caret)

## Loading required package: ggplot2

## Loading required package: lattice

Carseats_Filtered <- Carseats %>% select("Sales", "Price",
"Advertising", "Population", "Age", "Income", "Education")

set.seed(123)
trainIndex <- createDataPartition(Carseats_Filtered$Sales, p = 0.7, list =
FALSE)
Train_Data <- Carseats_Filtered[trainIndex, ]
Test_Data <- Carseats_Filtered[-trainIndex, ]

#Set up the model.
Support_Vector_Machine_Model <- train(Sales ~.,
data = Train_Data,
method = "svmLinear",
trControl = trainControl(method = "cv", number
= 10))

#Presenting the model's summary information.
summary(Support_Vector_Machine_Model)

## Length Class Mode
##      1  ksvm    S4

#Make predictions using the test dataset.
Predictions <- predict(Support_Vector_Machine_Model, newdata = Test_Data)

#Determine the R-squared value.
R_squared<- postResample(Predictions, Test_Data$Sales)
R_squared

##      RMSE Rsquared      MAE
## 2.297064 0.385761 1.876610

#QB2. Customize the search grid by checking the model's performance for C
parameter of 0.1,.5,1 and 10 using 2 repeats of 5-fold cross validation.

library(caret)
Grid <- expand.grid(C = c(0.1,0.5,1,10))
trctrl2 <- trainControl(method = "repeatedcv", number = 5, repeats = 2)
svm_Linear_Grid <- train(Sales~., data = Carseats_Filtered, method =
"svmLinear",
trControl=trctrl2,
preProcess = c("center", "scale"),
tuneGrid = Grid,

```

```

                                tuneLength = 10)
svm_Linear_Grid

## Support Vector Machines with Linear Kernel
##
## 400 samples
## 6 predictor
##
## Pre-processing: centered (6), scaled (6)
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 319, 320, 320, 321, 320, 320, ...
## Resampling results across tuning parameters:
##
##  C      RMSE      Rsquared  MAE
##  0.1  2.270125  0.3649327  1.818221
##  0.5  2.269631  0.3642787  1.816916
##  1.0  2.269468  0.3642989  1.816430
## 10.0  2.269407  0.3643307  1.816397
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was C = 10.

#QB3. Train a neural network model to predict Sales based on all other
attributes ("Price", "Advertising", "Population", "Age", "Income" and
"Education"). Hint: use caret train() with method set to "nnet". What is the
R-square of the model with the best hyper parameters (using default caret
search grid) - hint: don't forget to scale the data.

#Data scaling
scaled <- preProcess(Carseats_Filtered[-1], method = "scale")
train <- predict(scaled, Carseats_Filtered[-1])
#Train control
set.seed(27)
folds <- trainControl(method = "repeatedcv",
                      number = 5,
                      repeats = 2,
                      verboseIter = FALSE)
# Using default search grid, training the neural network model
set.seed(123)
nnet_Cars <- train(Sales~., data = Carseats_Filtered ,
                  method = "nnet",
                  trControl = folds)

## # weights: 9
## initial value 17583.375508
## final value 15744.713000
## converged
## # weights: 25
## initial value 18250.380996

```

```
## final value 15744.713000
## converged
## # weights: 41
## initial value 17969.872043
## final value 15744.713000
## converged
## # weights: 9
## initial value 17485.862634
## iter 10 value 15752.495629
## iter 20 value 15749.181183
## final value 15749.165782
## converged
## # weights: 25
## initial value 16731.281643
## iter 10 value 15750.726329
## iter 20 value 15747.342759
## iter 30 value 15747.245743
## iter 30 value 15747.245604
## iter 30 value 15747.245554
## final value 15747.245554
## converged
## # weights: 41
## initial value 17717.818624
## iter 10 value 15761.704139
## iter 20 value 15746.593674
## final value 15746.528022
## converged
## # weights: 9
## initial value 18907.374924
## iter 10 value 15751.310870
## iter 20 value 15744.789068
## iter 20 value 15744.789038
## final value 15744.789038
## converged
## # weights: 25
## initial value 17691.631949
## iter 10 value 15766.322138
## iter 20 value 15744.962136
## iter 30 value 15744.727758
## iter 30 value 15744.727752
## iter 30 value 15744.727746
## final value 15744.727746
## converged
## # weights: 41
## initial value 18730.502990
## iter 10 value 15752.934200
## iter 20 value 15744.986373
## final value 15744.735157
## converged
## # weights: 9
```

```
## initial value 18150.262633
## final value 16133.992900
## converged
## # weights: 25
## initial value 18057.338235
## final value 16133.992900
## converged
## # weights: 41
## initial value 17359.482931
## final value 16133.992900
## converged
## # weights: 9
## initial value 17989.498119
## iter 10 value 16200.459178
## iter 20 value 16138.492102
## final value 16138.454408
## converged
## # weights: 25
## initial value 18745.698886
## iter 10 value 16153.006756
## iter 20 value 16137.256703
## iter 30 value 16136.548298
## final value 16136.530460
## converged
## # weights: 41
## initial value 17459.477458
## iter 10 value 16143.635989
## iter 20 value 16136.062684
## iter 30 value 16135.826166
## iter 40 value 16135.818629
## iter 40 value 16135.818513
## final value 16135.811494
## converged
## # weights: 9
## initial value 17665.772198
## iter 10 value 16142.809710
## iter 20 value 16134.094551
## iter 20 value 16134.094510
## final value 16134.094510
## converged
## # weights: 25
## initial value 17838.240242
## iter 10 value 16143.041882
## iter 20 value 16134.097228
## iter 30 value 16134.005971
## iter 30 value 16134.005867
## iter 30 value 16134.005862
## final value 16134.005862
## converged
## # weights: 41
```

```
## initial value 17513.611852
## iter 10 value 16144.718662
## iter 20 value 16134.116560
## final value 16134.008897
## converged
## # weights: 9
## initial value 18527.696365
## final value 16054.878300
## converged
## # weights: 25
## initial value 17788.524754
## final value 16054.878300
## converged
## # weights: 41
## initial value 17446.488348
## final value 16054.878300
## converged
## # weights: 9
## initial value 18790.788824
## iter 10 value 16060.229113
## iter 20 value 16059.341696
## final value 16059.340150
## converged
## # weights: 25
## initial value 18626.938869
## iter 10 value 16082.809489
## iter 20 value 16057.679412
## iter 30 value 16057.431616
## final value 16057.416615
## converged
## # weights: 41
## initial value 18338.133205
## iter 10 value 16153.067658
## iter 20 value 16058.367215
## iter 30 value 16057.518420
## iter 40 value 16056.957883
## iter 50 value 16056.696790
## iter 50 value 16056.696704
## iter 50 value 16056.696682
## final value 16056.696682
## converged
## # weights: 9
## initial value 17590.373920
## iter 10 value 16059.385823
## iter 20 value 16054.930268
## iter 20 value 16054.930247
## final value 16054.930247
## converged
## # weights: 25
## initial value 17195.152150
```

```
## iter 10 value 16063.994339
## iter 20 value 16054.983401
## final value 16054.887538
## converged
## # weights: 41
## initial value 19381.816748
## iter 10 value 16074.408053
## iter 20 value 16055.103463
## iter 30 value 16054.916332
## final value 16054.893441
## converged
## # weights: 9
## initial value 18955.393490
## final value 16058.233500
## converged
## # weights: 25
## initial value 18731.157319
## final value 16058.233500
## converged
## # weights: 41
## initial value 17519.676097
## final value 16058.233500
## converged
## # weights: 9
## initial value 18443.945205
## iter 10 value 16062.703228
## final value 16062.691766
## converged
## # weights: 25
## initial value 17324.504702
## iter 10 value 16063.061711
## iter 20 value 16060.975591
## iter 30 value 16060.769179
## iter 30 value 16060.769058
## iter 30 value 16060.768932
## final value 16060.768932
## converged
## # weights: 41
## initial value 17182.149602
## iter 10 value 16062.517552
## iter 20 value 16060.350351
## iter 30 value 16060.151610
## iter 40 value 16060.063571
## final value 16060.050628
## converged
## # weights: 9
## initial value 17573.438755
## iter 10 value 16062.663964
## iter 20 value 16058.284580
## iter 20 value 16058.284559
```



```
## final value 16058.284559
## converged
## # weights: 25
## initial value 17709.202204
## iter 10 value 16070.257094
## iter 20 value 16058.372123
## final value 16058.252642
## converged
## # weights: 41
## initial value 18558.603853
## iter 10 value 16075.463049
## iter 20 value 16058.432143
## final value 16058.249622
## converged
## # weights: 9
## initial value 18722.292207
## final value 16260.862700
## converged
## # weights: 25
## initial value 17799.426161
## final value 16260.862700
## converged
## # weights: 41
## initial value 18255.742928
## final value 16260.862700
## converged
## # weights: 9
## initial value 17481.607340
## iter 10 value 16273.839719
## iter 20 value 16265.331018
## final value 16265.323262
## converged
## # weights: 25
## initial value 18295.365403
## iter 10 value 16326.205795
## iter 20 value 16265.170871
## iter 30 value 16263.576040
## final value 16263.399405
## converged
## # weights: 41
## initial value 17990.022132
## iter 10 value 16263.841721
## iter 20 value 16262.695365
## final value 16262.680613
## converged
## # weights: 9
## initial value 18295.931560
## iter 10 value 16271.756840
## iter 20 value 16260.988301
## iter 20 value 16260.988251
```

```
## final value 16260.988251
## converged
## # weights: 25
## initial value 18603.949520
## iter 10 value 16276.382456
## iter 20 value 16261.041631
## iter 30 value 16260.871843
## iter 30 value 16260.871810
## final value 16260.870307
## converged
## # weights: 41
## initial value 18012.253619
## iter 10 value 16277.963842
## iter 20 value 16261.059863
## iter 30 value 16260.871990
## iter 30 value 16260.871973
## final value 16260.871512
## converged
## # weights: 9
## initial value 17699.831901
## final value 16275.959300
## converged
## # weights: 25
## initial value 19641.795202
## final value 16275.959300
## converged
## # weights: 41
## initial value 18752.645654
## final value 16275.959300
## converged
## # weights: 9
## initial value 17240.269435
## iter 10 value 16283.782698
## final value 16280.421142
## converged
## # weights: 25
## initial value 18553.272473
## iter 10 value 16279.424338
## iter 20 value 16278.558130
## iter 30 value 16278.527772
## final value 16278.496768
## converged
## # weights: 41
## initial value 17756.873961
## iter 10 value 16280.244531
## iter 20 value 16277.829378
## iter 30 value 16277.780248
## final value 16277.777787
## converged
## # weights: 9
```

```
## initial value 19121.132256
## iter 10 value 16284.843993
## iter 20 value 16276.061734
## iter 20 value 16276.061693
## final value 16276.061693
## converged
## # weights: 25
## initial value 18036.303840
## iter 10 value 16281.085424
## iter 20 value 16276.018400
## final value 16275.980929
## converged
## # weights: 41
## initial value 19355.400213
## iter 10 value 16288.718731
## iter 20 value 16276.106406
## final value 16275.970733
## converged
## # weights: 9
## initial value 18050.702295
## final value 15928.817500
## converged
## # weights: 25
## initial value 18673.513835
## final value 15928.817500
## converged
## # weights: 41
## initial value 18478.275519
## final value 15928.817500
## converged
## # weights: 9
## initial value 18032.649747
## iter 10 value 15935.432118
## final value 15933.274268
## converged
## # weights: 25
## initial value 18953.220557
## iter 10 value 15937.188701
## iter 20 value 15932.025175
## iter 30 value 15931.704036
## iter 40 value 15931.355821
## final value 15931.352118
## converged
## # weights: 41
## initial value 17775.320098
## iter 10 value 15931.261131
## iter 20 value 15930.730135
## iter 30 value 15930.634420
## iter 30 value 15930.634354
## iter 30 value 15930.634354
```

```
## final value 15930.634354
## converged
## # weights: 9
## initial value 17439.224458
## iter 10 value 15933.303079
## iter 20 value 15928.869215
## iter 20 value 15928.869195
## final value 15928.869195
## converged
## # weights: 25
## initial value 17831.381529
## iter 10 value 15939.629190
## iter 20 value 15928.942150
## final value 15928.827000
## converged
## # weights: 41
## initial value 16885.798107
## iter 10 value 15943.011273
## iter 20 value 15928.981143
## final value 15928.827662
## converged
## # weights: 9
## initial value 19001.093244
## final value 16026.916900
## converged
## # weights: 25
## initial value 18234.740751
## final value 16026.916900
## converged
## # weights: 41
## initial value 17444.773107
## final value 16026.916900
## converged
## # weights: 9
## initial value 19226.609443
## iter 10 value 16031.661108
## final value 16031.376945
## converged
## # weights: 25
## initial value 17912.050438
## iter 10 value 16099.176152
## iter 20 value 16031.084078
## iter 30 value 16030.127163
## final value 16030.125564
## converged
## # weights: 41
## initial value 19636.734646
## iter 10 value 16041.491939
## iter 20 value 16030.108778
## iter 30 value 16028.945067
```

```
## iter 40 value 16028.760813
## iter 50 value 16028.737686
## final value 16028.735586
## converged
## # weights: 9
## initial value 17894.055196
## iter 10 value 16037.383250
## iter 20 value 16027.037569
## iter 20 value 16027.037521
## final value 16027.037521
## converged
## # weights: 25
## initial value 17695.042357
## iter 10 value 16036.926838
## iter 20 value 16027.032307
## final value 16026.947386
## converged
## # weights: 41
## initial value 16889.040753
## iter 10 value 16033.782658
## iter 20 value 16026.996057
## final value 16026.924364
## converged
## # weights: 9
## initial value 18092.146826
## final value 16199.115500
## converged
## # weights: 25
## initial value 18173.879067
## final value 16199.115500
## converged
## # weights: 41
## initial value 19348.973466
## final value 16199.115500
## converged
## # weights: 9
## initial value 17686.176332
## iter 10 value 16205.168079
## final value 16203.578217
## converged
## # weights: 25
## initial value 18221.044579
## iter 10 value 16206.649650
## iter 20 value 16201.780743
## iter 30 value 16201.659845
## final value 16201.653530
## converged
## # weights: 41
## initial value 18007.644613
## iter 10 value 16236.916646
```

```
## iter 20 value 16201.368158
## iter 30 value 16200.961996
## final value 16200.935019
## converged
## # weights: 9
## initial value 18647.736276
## iter 10 value 16206.810589
## iter 20 value 16199.204218
## iter 20 value 16199.204183
## final value 16199.204183
## converged
## # weights: 25
## initial value 17681.544301
## iter 10 value 16203.519924
## iter 20 value 16199.166280
## final value 16199.138361
## converged
## # weights: 41
## initial value 18978.786685
## iter 10 value 16214.630998
## iter 20 value 16199.294382
## iter 30 value 16199.129205
## iter 30 value 16199.129191
## final value 16199.128531
## converged
## # weights: 9
## initial value 18177.074240
## final value 15821.871200
## converged
## # weights: 25
## initial value 18778.113264
## final value 15821.871200
## converged
## # weights: 41
## initial value 16909.381852
## final value 15821.871200
## converged
## # weights: 9
## initial value 17119.260930
## iter 10 value 15826.599396
## iter 20 value 15826.324908
## final value 15826.324509
## converged
## # weights: 25
## initial value 18703.665234
## iter 10 value 15832.759440
## iter 20 value 15826.738148
## iter 30 value 15824.517679
## iter 40 value 15824.417539
## iter 50 value 15824.405208
```

```
## final value 15824.404153
## converged
## # weights: 41
## initial value 18070.771384
## iter 10 value 15824.511212
## iter 20 value 15823.959776
## iter 30 value 15823.687857
## final value 15823.686523
## converged
## # weights: 9
## initial value 18880.142624
## iter 10 value 15828.619435
## iter 20 value 15821.949002
## iter 20 value 15821.948971
## final value 15821.948971
## converged
## # weights: 25
## initial value 17263.346538
## iter 10 value 15823.989488
## iter 20 value 15821.994572
## final value 15821.983647
## converged
## # weights: 41
## initial value 17737.097268
## iter 10 value 15841.716075
## iter 20 value 15822.099996
## final value 15821.883136
## converged

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info =
trainInfo,
## : There were missing values in resampled performance measures.

## # weights: 9
## initial value 23385.579068
## iter 10 value 20079.130735
## iter 20 value 20063.354114
## iter 20 value 20063.354040
## final value 20063.354040
## converged

nnet_Cars

## Neural Network
##
## 400 samples
## 6 predictor
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 320, 321, 319, 320, 320, 319, ...
```

```
## Resampling results across tuning parameters:
```

```
##
```

##	size	decay	RMSE	Rsquared	MAE
##	1	0e+00	7.080831	NaN	6.511777
##	1	1e-04	7.080831	NaN	6.511777
##	1	1e-01	7.081018	0.02150646	6.511976
##	3	0e+00	7.080831	NaN	6.511777
##	3	1e-04	7.080831	0.05189222	6.511777
##	3	1e-01	7.080934	0.03272925	6.511887
##	5	0e+00	7.080831	NaN	6.511777
##	5	1e-04	7.080831	0.03626897	6.511777
##	5	1e-01	7.080900	0.02608450	6.511851

```
##
```

```
## RMSE was used to select the optimal model using the smallest value.
```

```
## The final values used for the model were size = 1 and decay = 1e-04.
```

#The optimal model was chosen based on the RMSE (root mean squared error), featuring a size of 1 and a decay of 0. The RMSE for this model was 7.081637, accompanied by a corresponding MAE (mean absolute error) value of 6.511536. While the Rsquared value was 'not available' for the optimal model, it ranged from NaN to 0.02538470 for other models.

#QB4 - "Consider the following input: Sales=9, Price=6.54, Population=124, Advertising=0, Age=76, Income= 110, Education=10 What will be the estimated Sales for this record using the above neuralnet model?"

```
Sales <- c(9)
```

```
Price <- c(6.54)
```

```
Population <- c(124)
```

```
Advertising <- c(0)
```

```
Age <- c(76)
```

```
Income <- c(110)
```

```
Education <- c(10)
```

```
Test <- data.frame(Sales, Price, Population, Advertising, Age, Income, Education)
```

```
# Making Predictions
```

```
Predicting_sales <- predict(nnet_Cars, Test)
```

```
Predicting_sales
```

```
## 1
```

```
## 1
```